

# High-Performance Development Environment Setup

*Optimized for ThinkPad P1 Gen 7 - 64GB RAM, 22-core Intel Ultra 7 165H*

## Phase 1: Foundation Setup

### 1. Development Tools Installation

bash

*# Essential development packages*

```
sudo dnf groupinstall "Development Tools" "Development Libraries"
```

```
sudo dnf install -y \  
    git git-lfs \  
    neovim code \  
    docker podman \  
    kubernetes kubectl helm \  
    nodejs npm yarn \  
    python3 python3-pip python3-venv \  
    go lang rust cargo \  
    java-17-openjdk maven gradle \  
    postgresql-server redis \  
    nginx \  
    tmux zellij \  
    btop htop \  
    jq yq \  
    terraform \  
    ansible
```

*# Modern CLI tools for your .zshrc*

```
sudo dnf install -y \  
    bat exa ripgrep fd-find \  
   procs dust duf \  
    git-delta fzf \  
    zoxide starship
```

### 2. Container Platform Setup (Leverage that 64GB RAM!)

```
bash
```

```
# Configure Podman for rootless containers
```

```
echo 'unqualified-search-registries = ["docker.io"]' | sudo tee -a /etc/containers/regi
```

```
# Increase container limits for your massive RAM
```

```
mkdir -p ~/.config/containers
```

```
cat > ~/.config/containers/storage.conf << 'EOF'
```

```
[storage]
```

```
driver = "overlay"
```

```
runroot = "/run/user/1000/containers"
```

```
graphroot = "/home/$USER/.local/share/containers/storage"
```

```
[storage.options]
```

```
size = "100G"
```

```
EOF
```

```
# Set resource limits to use your full potential
```

```
echo "vm.max_map_count=262144" | sudo tee -a /etc/sysctl.conf
```

```
sudo sysctl -w vm.max_map_count=262144
```



## 🎯 Phase 2: Multi-Environment Development Setup

### 1. Local Kubernetes Development

```
bash
```

```
# Install K3d for local Kubernetes (lightweight, perfect for development)
```

```
curl -s https://raw.githubusercontent.com/k3d-io/k3d/main/install.sh | bash
```

```
# Create a powerful local cluster using your resources
```

```
k3d cluster create dev-cluster \
```

```
--agents 3 \
```

```
--memory 16g \
```

```
--cpus 8 \
```

```
--api-port 6443 \
```

```
--port 8080:80@loadbalancer \
```

```
--port 8443:443@loadbalancer
```

```
# Install Kubernetes dashboard
```

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.7.0/aio/depl
```



## 2. Development Databases (Containerized)

bash

*# PostgreSQL for serious applications*

```
podman run -d \  
  --name postgres-dev \  
  -e POSTGRES_PASSWORD=devpass \  
  -e POSTGRES_DB=devdb \  
  -p 5432:5432 \  
  -v postgres-data:/var/lib/postgresql/data \  
  postgres:15
```

*# Redis for caching/sessions*

```
podman run -d \  
  --name redis-dev \  
  -p 6379:6379 \  
  redis:7-alpine
```

*# MongoDB for NoSQL projects*

```
podman run -d \  
  --name mongo-dev \  
  -e MONGO_INITDB_ROOT_USERNAME=admin \  
  -e MONGO_INITDB_ROOT_PASSWORD=devpass \  
  -p 27017:27017 \  
  -v mongo-data:/data/db \  
  mongo:6
```

*# ClickHouse for analytics (your RAM can handle it!)*

```
podman run -d \  
  --name clickhouse-dev \  
  -p 8123:8123 \  
  -p 9000:9000 \  
  --ulimit nofile=262144:262144 \  
  -v clickhouse-data:/var/lib/clickhouse \  
  clickhouse/clickhouse-server:latest
```

## 3. Message Queues & Streaming

bash

*# Apache Kafka (memory-intensive, perfect for your setup)*

```
podman run -d \  
  --name zookeeper \  
  -p 2181:2181 \  
  -e ZOOKEEPER_CLIENT_PORT=2181 \  
  confluentinc/cp-zookeeper:latest
```

```
podman run -d \  
  --name kafka \  
  -p 9092:9092 \  
  -e KAFKA_ZOOKEEPER_CONNECT=localhost:2181 \  
  -e KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://localhost:9092 \  
  -e KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR=1 \  
  -e KAFKA_HEAP_OPTS="-Xmx2G -Xms2G" \  
  confluentinc/cp-kafka:latest
```

*# RabbitMQ for traditional messaging*

```
podman run -d \  
  --name rabbitmq-dev \  
  -p 5672:5672 \  
  -p 15672:15672 \  
  -e RABBITMQ_DEFAULT_USER=admin \  
  -e RABBITMQ_DEFAULT_PASS=devpass \  
  rabbitmq:3-management
```

## Phase 3: AI/ML Development Environment

### 1. Python AI/ML Stack

```
bash
```

```
# Create dedicated Python environment for ML
```

```
python3 -m venv ~/venv/ml
```

```
source ~/venv/ml/bin/activate
```

```
# Install core ML libraries (your 64GB RAM can handle large models!)
```

```
pip install \
```

```
    torch torchvision torchaudio \
```

```
    tensorflow \
```

```
    scikit-learn pandas numpy \
```

```
    jupyter jupyterlab \
```

```
    transformers datasets \
```

```
    langchain openai \
```

```
    streamlit \
```

```
    mlflow wandb \
```

```
    dvc
```

```
# Start Jupyter Lab with increased resources
```

```
echo "c.NotebookApp.max_buffer_size = 2**28" > ~/.jupyter/jupyter_notebook_config.py
```

## 2. Local AI Model Serving

```
bash
```

```
# Ollama for local LLM serving (leveraging your hardware)
```

```
curl -fsSL https://ollama.com/install.sh | sh
```

```
# Start Ollama and pull models
```

```
systemctl --user enable ollama
```

```
systemctl --user start ollama
```

```
# Pull some models (your RAM can handle multiple large models)
```

```
ollama pull llama2:13b
```

```
ollama pull codellama:34b
```

```
ollama pull mistral:7b
```

## Phase 4: Development Tool Configuration

### 1. VS Code with Powerhouse Extensions

bash

```
# Install VS Code extensions for multi-language development
code --install-extension ms-python.python
code --install-extension golang.go
code --install-extension rust-lang.rust-analyzer
code --install-extension ms-vscode.vscode-typescript-next
code --install-extension ms-kubernetes-tools.vscode-kubernetes-tools
code --install-extension ms-azuretools.vscode-docker
code --install-extension hashicorp.terraform
code --install-extension redhat.ansible
code --install-extension ms-toolsai.jupyter
code --install-extension github.copilot

# Configure VS Code for your hardware
mkdir -p ~/.config/Code/User
cat > ~/.config/Code/User/settings.json << 'EOF'
{
  "files.watcherExclude": {
    "**/node_modules/**": true
  },
  "search.maxResults": 50000,
  "terminal.integrated.profiles.linux": {
    "zsh": {
      "path": "/usr/bin/zsh"
    }
  },
  "terminal.integrated.defaultProfile.linux": "zsh",
  "python.defaultInterpreterPath": "~/venv/ml/bin/python",
  "go.toolsManagement.autoUpdate": true,
  "rust-analyzer.server.path": "rust-analyzer",
  "kubernetes.kubectl-path.linux": "/usr/bin/kubectl"
}
EOF
```

## 2. Neovim Configuration (For Terminal Power Users)

bash

*# Install LazyVim (modern Neovim distribution)*

```
git clone https://github.com/LazyVim/starter ~/.config/nvim
```

```
rm -rf ~/.config/nvim/.git
```

*# Install language servers for multi-language support*

```
sudo dnf install -y \
```

```
    nodejs npm \
```

```
    python3-pip \
```

```
    rust-analyzer \
```

```
    gopls
```



## Phase 5: Workflow Automation

### 1. Development Environment Manager

bash

*# Create environment control script*

cat > ~/bin/dev-env << 'EOF'

#!/bin/bash

case \$1 in

start)

echo "🚀 Starting development environment..."

systemctl --user start podman

podman start postgres-dev redis-dev mongo-dev clickhouse-dev rabbitmq-dev 2>/dev/null

k3d cluster start dev-cluster 2>/dev/null

systemctl --user start ollama

echo "✅ Development environment ready!"

;;

stop)

echo "🛑 Stopping development environment..."

podman stop postgres-dev redis-dev mongo-dev clickhouse-dev rabbitmq-dev 2>/dev/null

k3d cluster stop dev-cluster 2>/dev/null

systemctl --user stop ollama

echo "✅ Development environment stopped!"

;;

status)

echo "📊 Development Environment Status:"

echo "Containers:"

podman ps --format "table {{.Names}}\t{{.Status}}\t{{.Ports}}"

echo -e "\nKubernetes:"

k3d cluster list

echo -e "\nSystem Resources:"

free -h

;;

restart)

\$0 stop

sleep 5

\$0 start

;;

\*)

echo "Usage: dev-env {start|stop|status|restart}"

;;

esac

EOF

chmod +x ~/bin/dev-env



## 2. Resource Monitoring Dashboard

```
bash
```

```
# Install modern system monitoring  
sudo dnf install -y btop bandwhich
```

```
# Create monitoring script
```

```
cat > ~/bin/monitor << 'EOF'
```

```
#!/bin/bash
```

```
# Multi-pane monitoring for development
```

```
tmux new-session -d -s monitor
```

```
tmux split-window -h
```

```
tmux split-window -v
```

```
tmux select-pane -t 0
```

```
tmux split-window -v
```

```
tmux send-keys -t 0 'btop' Enter
```

```
tmux send-keys -t 1 'watch -n 2 "podman ps --format \"table {{.Names}}\t{{.Status}}\t{'
```

```
tmux send-keys -t 2 'watch -n 5 "kubectl get pods -A"' Enter
```

```
tmux send-keys -t 3 'journalctl -f -u podman' Enter
```

```
tmux attach-session -t monitor
```

```
EOF
```

```
chmod +x ~/bin/monitor
```

## Phase 6: Language-Specific Optimizations

### 1. Go Development (Optimized for Fast Compilation)

```
bash
```

```
# Configure Go for your multicore CPU
```

```
echo 'export GOMAXPROCS=22' >> ~/.zshrc
```

```
echo 'export GOPATH=~/.cache/go-build' >> ~/.zshrc
```

```
# Install Go tools
```

```
go install golang.org/x/tools/gopls@latest
```

```
go install github.com/golangci/golangci-lint/cmd/golangci-lint@latest
```

```
go install github.com/air-verse/air@latest
```

## 2. Rust Development (Leverage Parallel Compilation)

```
bash

# Configure Rust for parallel builds
echo 'export CARGO_BUILD_JOBS=22' >> ~/.zshrc
echo 'export RUSTC_WRAPPER=sccache' >> ~/.zshrc

# Install sccache for compilation caching
cargo install sccache
```

## 3. Node.js Development (Multiple Versions)

```
bash

# Install Node Version Manager
curl -o- https://raw.githubusercontent.com/nvm-sh/nvm/v0.39.0/install.sh | bash

# Install multiple Node versions for different projects
nvm install --lts
nvm install 18
nvm install 20
nvm use --lts

# Configure npm for your setup
npm config set cache ~/.npm-cache
npm config set maxsockets 50
```

## Phase 7: Performance Testing Environment

### 1. Load Testing Tools

```
bash

# Install performance testing tools
sudo dnf install -y apache-bench siege
cargo install drill
go install github.com/tsenart/vegeta@latest

# Install k6 for modern load testing
sudo dnf install xz
curl -s https://github.com/grafana/k6/releases/download/v0.47.0/k6-v0.47.0-linux-amd64.
sudo mv k6-v0.47.0-linux-amd64/k6 /usr/local/bin/
```

## 2. Monitoring Stack (Prometheus + Grafana)

bash

*# Prometheus for metrics*

```
podman run -d \  
  --name prometheus \  
  -p 9090:9090 \  
  -v prometheus-data:/prometheus \  
  prom/prometheus:latest
```

*# Grafana for visualization*

```
podman run -d \  
  --name grafana \  
  -p 3000:3000 \  
  -e GF_SECURITY_ADMIN_PASSWORD=admin \  
  -v grafana-data:/var/lib/grafana \  
  grafana/grafana:latest
```

## Phase 8: Quick Start Commands

Add these to your `.zshrc`:

bash

*# Development environment shortcuts*

alias dev-start='~/bin/dev-env start'

alias dev-stop='~/bin/dev-env stop'

alias dev-status='~/bin/dev-env status'

alias dev-monitor='~/bin/monitor'

*# Quick service connections*

alias pg-connect='psql postgresql://postgres:devpass@localhost:5432/devdb'

alias redis-cli='redis-cli -h localhost -p 6379'

alias mongo-connect='mongosh "mongodb://admin:devpass@localhost:27017/admin"'

*# Container management*

alias containers='podman ps --format "table {{.Names}}\t{{.Status}}\t{{.Ports}}"'

alias container-stats='podman stats'

*# Kubernetes shortcuts*

alias k='kubectl'

alias pods='kubectl get pods -A'

alias services='kubectl get services -A'

*# Resource monitoring*

alias resources='btop'

alias network='bandwhich'

alias processes='procs'